

Sentiment Analysis of Tweets

STAT 362-0 Final Project

Jimmy Zhang



Problem and Motivation

- Natural Language Processing (NLP)
 - Multiclass Classification (13 classes)
- Applications
 - Social media moderation
 - Hate speech detection

Dataset Overview

- [Kaggle Dataset](#)
- 40,000 Observations
 - Tweet Text (feature)
 - Sentiment Label (target)
 - Tweet/User ID (unused)
- Word Embeddings with SentenceTransformer
 - Sentence Tokenization
 - Sequence Generation




Figure 1. Distribution of sentiment classes in the dataset (13 classes).

Deep Learning Approach

- Pre-trained Embedding Models
 - Complexity vs Computational Cost tradeoff
- Baseline Models
 - Multilayer Perceptron (MLP)
 - Gradient Boosted Tree (XGBoost)
- Recurrent Neural Networks (RNNs) for Sequence Processing
 - Vanilla RNN
 - Gated RNN (LSTM, GRU)
 - Bidirectional RNN
- Transformer Model
- Different Architectures (Regularization, Dropout, etc.)

Experimental Setup

- Training/Validation/Test Sets (60/20/20 ratio)
- Compare 4 variants per model type
 - Select the best from each for overall analysis
- Multiclass Evaluation Metrics
 - Loss Function: Cross-Entropy
 - Final Performance: Accuracy
- Implementation
 - Google Colab: T4 GPU (High-RAM)
 - Local VSCode Poetry Environment

Key Results

- Training/Validation Loss Curves
- Model Ranking on Testing Set
- Confusion Matrix




Figure 2. Final training and validation loss curves for the best variant of each model type.

| Model | Test Accuracy | Test Loss |
|-----------------|---------------|-----------|
| 0 Bidirectional | 0.368 | 1.824 |
| 1 LSTM | 0.364 | 1.865 |
| 2 MLP | 0.361 | 1.858 |
| 3 GRU | 0.359 | 1.863 |
| 4 Transformer | 0.359 | 1.885 |
| 5 XGBoost | 0.348 | N/A |
| 6 RNN | 0.216 | 2.145 |

Table 1. Testing set performance of best models.




Figure 3. Confusion matrix from Bidirectional RNN predictions on testing set.

Conclusions

- Inherently a difficult task
 - Many classes with imbalance
 - Relatively limited training data compared to the real world
- Complexity doesn't necessarily improve performance
 - Simpler architectures learned more quickly
 - Use Dropout/Regularization to reduce overfitting
- Deep learning approach was moderately successful
 - Small but clear improvement upon baseline

Reflection

- I expected the workflow with running code in Google Colab and exporting to VSCode to be tricky, but that part worked smoothly
- The biggest challenge was figuring out what different architectures I should try for each model type, since I wanted to make changes that would at least give me insightful results, if not performance improvements
- My original planned approach worked successfully overall
- Lessons:
 - In the experimental setup, sometimes I just need to pick an arbitrary starting point for my first model before making changes later
 - Maintaining an organized workflow is very important to avoid accidental confusion and obtaining non-reproducible results

AI Disclosure

- Gemini autocomplete in Google Colab
 - Improve efficiency when reproducing large chunks of code with only minor changes
 - I did not accept suggestions unless I knew exactly what that code would do and it matched with what I had previously done
 - If I ran into issues, I would carefully debug step-by-step
- VSCode GitHub Copilot
 - Generate Docstrings for some of the custom-defined functions

References

- Chen, T. and Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785-794. <http://dx.doi.org/10.1145/2939672.2939785>.
- Chollet, F., et al. (2015). Keras. [https://keras.io/.](https://keras.io/)
- Reimers, N. and Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *arXiv:1908.10084*. <https://arxiv.org/abs/1908.10084>.
- Valarmathi, B., et al. (2024). Sentiment Analysis of Covid-19 Twitter Data using Deep Learning Algorithm. *Procedia Computer Science*, 235, 3397-3407. <https://doi.org/10.1016/j.procs.2024.04.320>.